

analysis

May 25, 2023

1 Imports and Setup

```
[62]: import pandas as pd
import numpy as np
```

```
[63]: df = pd.read_csv('data/starcraft_player_data.csv')
```

```
[64]: df
```

```
[64]:
```

	GameID	LeagueIndex	Age	HoursPerWeek	TotalHours	APM	\
0	52	5	27	10	3000	143.7180	
1	55	5	23	10	5000	129.2322	
2	56	4	30	10	200	69.9612	
3	57	3	19	20	400	107.6016	
4	58	3	32	10	500	122.8908	
...	
3390	10089	8	?	?	?	259.6296	
3391	10090	8	?	?	?	314.6700	
3392	10092	8	?	?	?	299.4282	
3393	10094	8	?	?	?	375.8664	
3394	10095	8	?	?	?	348.3576	
	SelectByHotkeys	AssignToHotkeys	UniqueHotkeys	MinimapAttacks	\		
0	0.003515	0.000220	7	0.000110			
1	0.003304	0.000259	4	0.000294			
2	0.001101	0.000336	4	0.000294			
3	0.001034	0.000213	1	0.000053			
4	0.001136	0.000327	2	0.000000			
...			
3390	0.020425	0.000743	9	0.000621			
3391	0.028043	0.001157	10	0.000246			
3392	0.028341	0.000860	7	0.000338			
3393	0.036436	0.000594	5	0.000204			
3394	0.029855	0.000811	4	0.000224			
	MinimapRightClicks	NumberOfPACs	GapBetweenPACs	ActionLatency	\		
0	0.000392	0.004849	32.6677	40.8673			

1	0.000432	0.004307	32.9194	42.3454
2	0.000461	0.002926	44.6475	75.3548
3	0.000543	0.003783	29.2203	53.7352
4	0.001329	0.002368	22.6885	62.0813
...
3390	0.000146	0.004555	18.6059	42.8342
3391	0.001083	0.004259	14.3023	36.1156
3392	0.000169	0.004439	12.4028	39.5156
3393	0.000780	0.004346	11.6910	34.8547
3394	0.001315	0.005566	20.0537	33.5142

	ActionsInPAC	TotalMapExplored	WorkersMade	UniqueUnitsMade	\
0	4.7508	28	0.001397		6
1	4.8434	22	0.001193		5
2	4.0430	22	0.000745		6
3	4.9155	19	0.000426		7
4	9.3740	15	0.001174		4
...
3390	6.2754	46	0.000877		5
3391	7.1965	16	0.000788		4
3392	6.3979	19	0.001260		4
3393	7.9615	15	0.000613		6
3394	6.3719	27	0.001566		7

	ComplexUnitsMade	ComplexAbilitiesUsed
0	0.000000	0.000000
1	0.000000	0.000208
2	0.000000	0.000189
3	0.000000	0.000384
4	0.000000	0.000019
...
3390	0.000000	0.000000
3391	0.000000	0.000000
3392	0.000000	0.000000
3393	0.000000	0.000631
3394	0.000457	0.000895

[3395 rows x 20 columns]

2 Cleaning Data

```
[65]: df_without_question_mark = df[df != '?'].dropna()
string_cols = ['TotalHours', 'HoursPerWeek', 'Age']
for s in string_cols:
    df_without_question_mark[s] = df_without_question_mark[s].astype(int)
df_without_question_mark
```

[65]:

	GameID	LeagueIndex	Age	HoursPerWeek	TotalHours	APM	\
0	52	5	27	10	3000	143.7180	
1	55	5	23	10	5000	129.2322	
2	56	4	30	10	200	69.9612	
3	57	3	19	20	400	107.6016	
4	58	3	32	10	500	122.8908	
...	
3335	9261	4	20	8	400	158.1390	
3336	9264	5	16	56	1500	186.1320	
3337	9265	4	21	8	100	121.6992	
3338	9270	3	20	28	400	134.2848	
3339	9271	4	22	6	400	88.8246	

	SelectByHotkeys	AssignToHotkeys	UniqueHotkeys	MinimapAttacks	\
0	0.003515	0.000220	7	0.000110	
1	0.003304	0.000259	4	0.000294	
2	0.001101	0.000336	4	0.000294	
3	0.001034	0.000213	1	0.000053	
4	0.001136	0.000327	2	0.000000	
...	
3335	0.013829	0.000504	7	0.000217	
3336	0.006951	0.000360	6	0.000083	
3337	0.002956	0.000241	8	0.000055	
3338	0.005424	0.000182	5	0.000000	
3339	0.000844	0.000108	2	0.000000	

	MinimapRightClicks	NumberOfPACs	GapBetweenPACs	ActionLatency	\
0	0.000392	0.004849	32.6677	40.8673	
1	0.000432	0.004307	32.9194	42.3454	
2	0.000461	0.002926	44.6475	75.3548	
3	0.000543	0.003783	29.2203	53.7352	
4	0.001329	0.002368	22.6885	62.0813	
...	
3335	0.000313	0.003583	36.3990	66.2718	
3336	0.000166	0.005414	22.8615	34.7417	
3337	0.000208	0.003690	35.5833	57.9585	
3338	0.000480	0.003205	18.2927	62.4615	
3339	0.000341	0.003099	45.1512	63.4435	

	ActionsInPAC	TotalMapExplored	WorkersMade	UniqueUnitsMade	\
0	4.7508	28	0.001397	6	
1	4.8434	22	0.001193	5	
2	4.0430	22	0.000745	6	
3	4.9155	19	0.000426	7	
4	9.3740	15	0.001174	4	
...	
3335	4.5097	30	0.001035	7	

3336	4.9309	38	0.001343	7
3337	5.4154	23	0.002014	7
3338	6.0202	18	0.000934	5
3339	5.1913	20	0.000476	8

	ComplexUnitsMade	ComplexAbilitiesUsed
0	0.0	0.000000
1	0.0	0.000208
2	0.0	0.000189
3	0.0	0.000384
4	0.0	0.000019
...
3335	0.0	0.000287
3336	0.0	0.000388
3337	0.0	0.000000
3338	0.0	0.000000
3339	0.0	0.000054

[3338 rows x 20 columns]

3 Exploring Correlations

```
[66]: corr = df_without_question_mark.corr()
```

```
[67]: corr.columns
```

```
[67]: Index(['GameID', 'LeagueIndex', 'Age', 'HoursPerWeek', 'TotalHours', 'APM',
          'SelectByHotkeys', 'AssignToHotkeys', 'UniqueHotkeys', 'MinimapAttacks',
          'MinimapRightClicks', 'NumberOfPACs', 'GapBetweenPACs', 'ActionLatency',
          'ActionsInPAC', 'TotalMapExplored', 'WorkersMade', 'UniqueUnitsMade',
          'ComplexUnitsMade', 'ComplexAbilitiesUsed'],
          dtype='object')
```

```
[68]: league_index_corr = corr['LeagueIndex']
```

```
[69]: league_index_corr
```

```
[69]: GameID          0.024974
      LeagueIndex     1.000000
      Age           -0.127518
      HoursPerWeek    0.217930
      TotalHours      0.023884
      APM             0.624171
      SelectByHotkeys 0.428637
      AssignToHotkeys 0.487280
      UniqueHotkeys   0.322415
```

```

MinimapAttacks      0.270526
MinimapRightClicks  0.206380
NumberOfPACs        0.589193
GapBetweenPACs      -0.537536
ActionLatency        -0.659940
ActionsInPAC         0.140303
TotalMapExplored     0.230347
WorkersMade          0.310452
UniqueUnitsMade      0.151933
ComplexUnitsMade     0.171190
ComplexAbilitiesUsed 0.156033
Name: LeagueIndex, dtype: float64

```

```
[70]: indices = np.where(abs(league_index_corr) > 0.5)
```

```
[71]: indices
```

```
[71]: (array([ 1,  5, 11, 12, 13]),)
```

```
[72]: corr.columns[indices]
```

```
[72]: Index(['LeagueIndex', 'APM', 'NumberOfPACs', 'GapBetweenPACs',
           'ActionLatency'],
           dtype='object')
```

From the **above** it appears that the categories with a relatively significant correlation to LeagueIndex (>0.5), which is our 1-8 code for rank, are **APM**, **NumberOfPACs**, **GapBetweenPACs**, and **ActionLatency**

So, it makes sense to continue by creating a classification model trained on this data. It is possible that adding in other features would only make our model worse because they are so uncorrelated with our output variable and could cause overfitting.

These variables make intuitive sense because APM, ActionLatency, NumberOfPACs, and GapBetweenPACs all correlate to how fast a player is and it makes sense that quicker players would have a higher rank because they have better reactions and more practice, developing their faster movement.

For further confirmation, I will also use scikit-learn's **SelectKBest** to see if selecting the 5 best features aligns with what we have above. I will use the `f_classif` because it is suitable for numerical classification.

```
[73]: features = list(df.columns)
      features.pop(1)
      print(features)
```

```

['GameID', 'Age', 'HoursPerWeek', 'TotalHours', 'APM', 'SelectByHotkeys',
'AssignToHotkeys', 'UniqueHotkeys', 'MinimapAttacks', 'MinimapRightClicks',
'NumberOfPACs', 'GapBetweenPACs', 'ActionLatency', 'ActionsInPAC',

```

```
'TotalMapExplored', 'WorkersMade', 'UniqueUnitsMade', 'ComplexUnitsMade',
'ComplexAbilitiesUsed']
```

```
[74]: from sklearn.feature_selection import SelectKBest, f_classif
X = df_without_question_mark[features]
y = df_without_question_mark['LeagueIndex']
kbest = SelectKBest(score_func=f_classif, k=5)
kbest.fit(X, y)
selected_features = kbest.get_support(indices=True)
print(df.columns[selected_features])
```

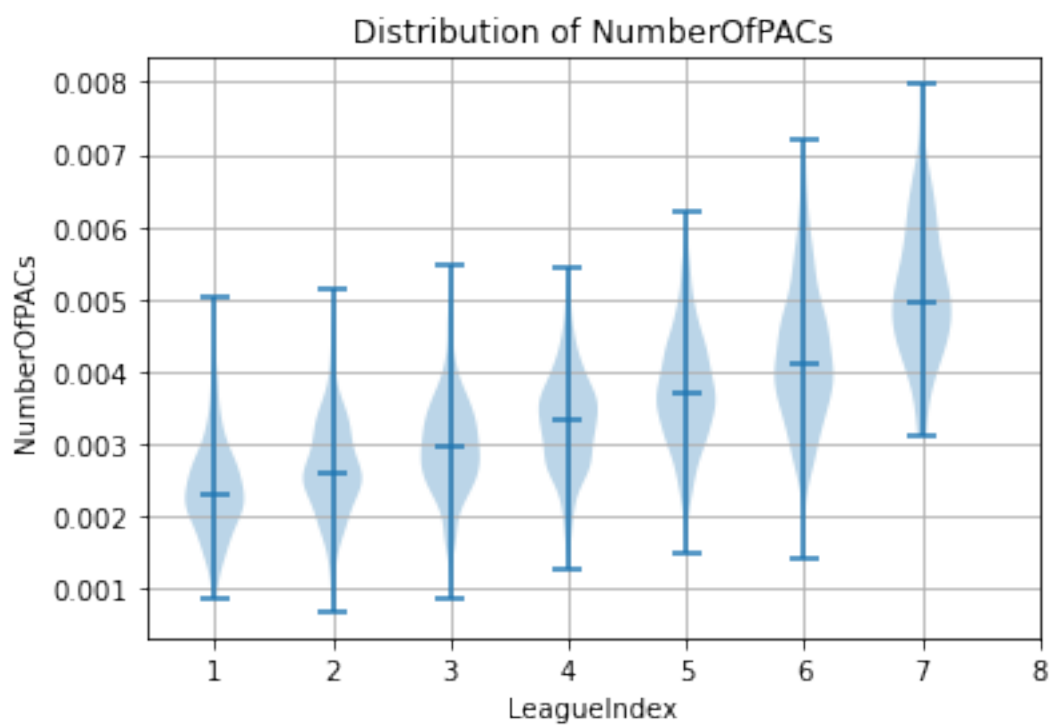
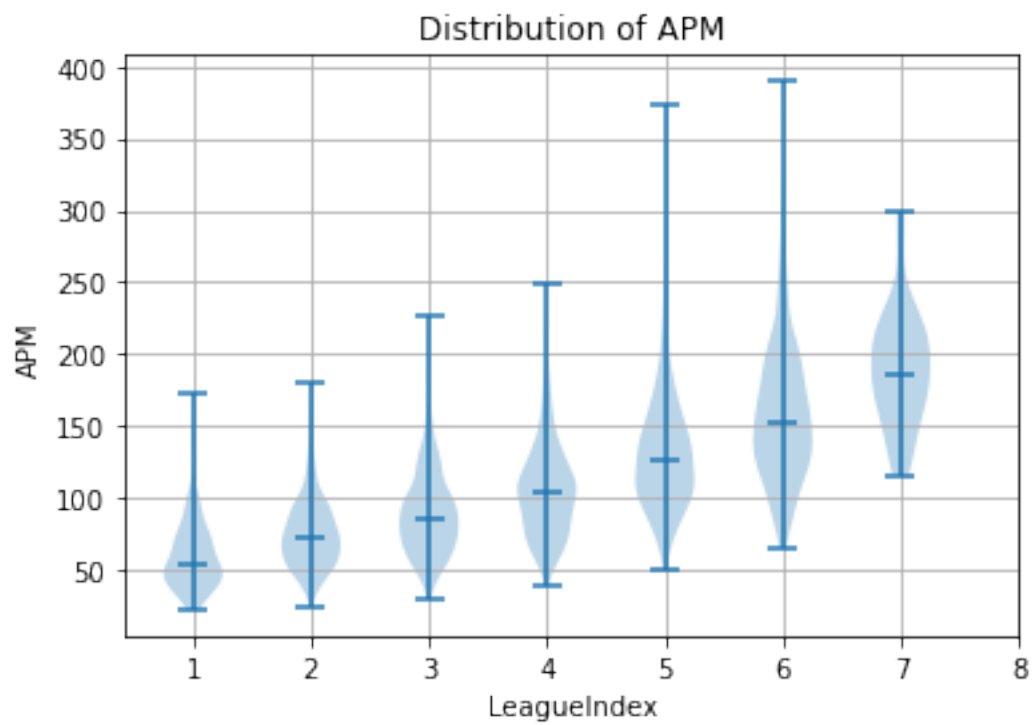
```
Index(['TotalHours', 'SelectByHotkeys', 'MinimapRightClicks', 'NumberOfPACs',
      'GapBetweenPACs'],
      dtype='object')
```

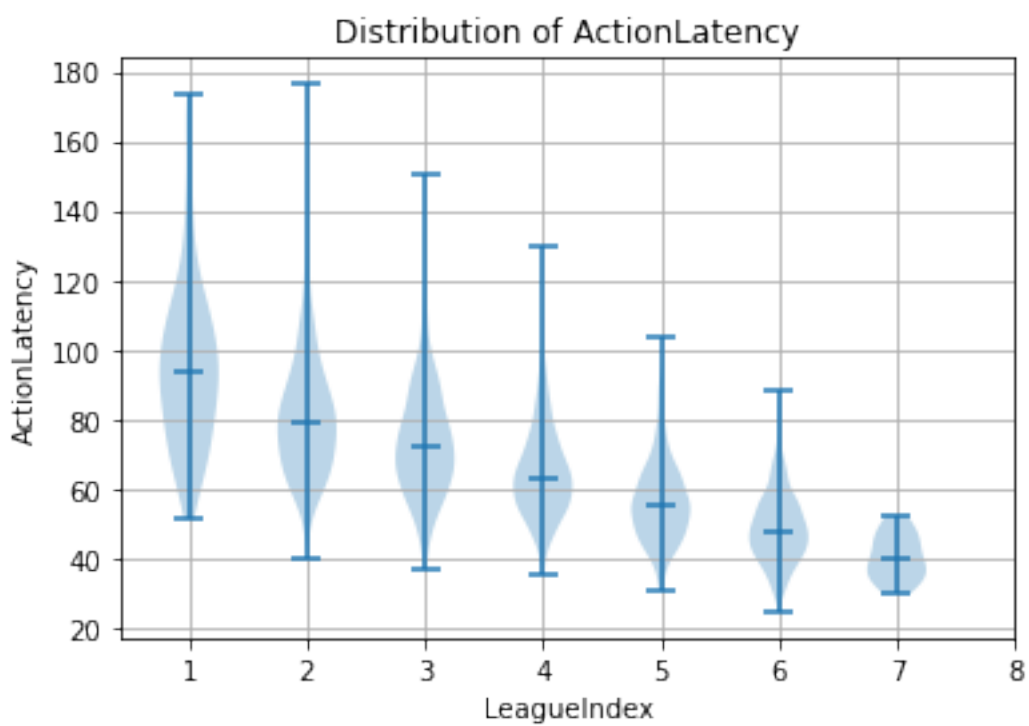
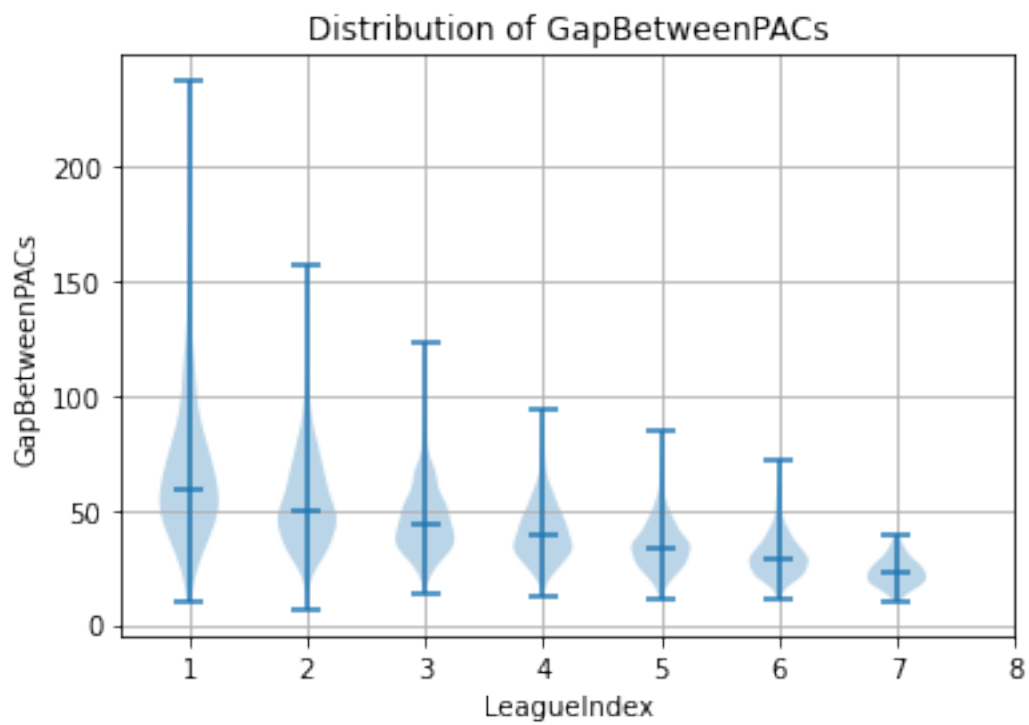
Performing this check we see that **TotalHours**, which was not included in the correlation matrix, is a good predictor for **LeagueIndex**, so this is an additional feature worth exploring. Scikit-Learn also suggests **MinimapRightClicks** is a good predictor, which makes sense because a more skilled player will check the minimap more often. Also, **SelectByHotkeys** measures some efficiency of game play, so also makes sense to correlate with skill

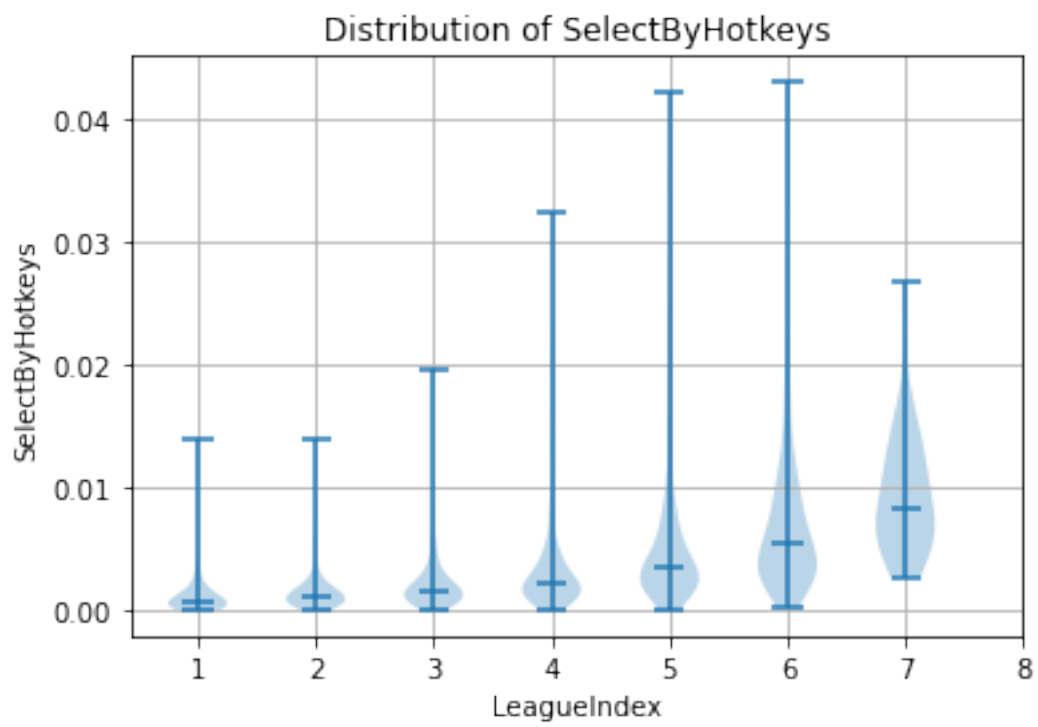
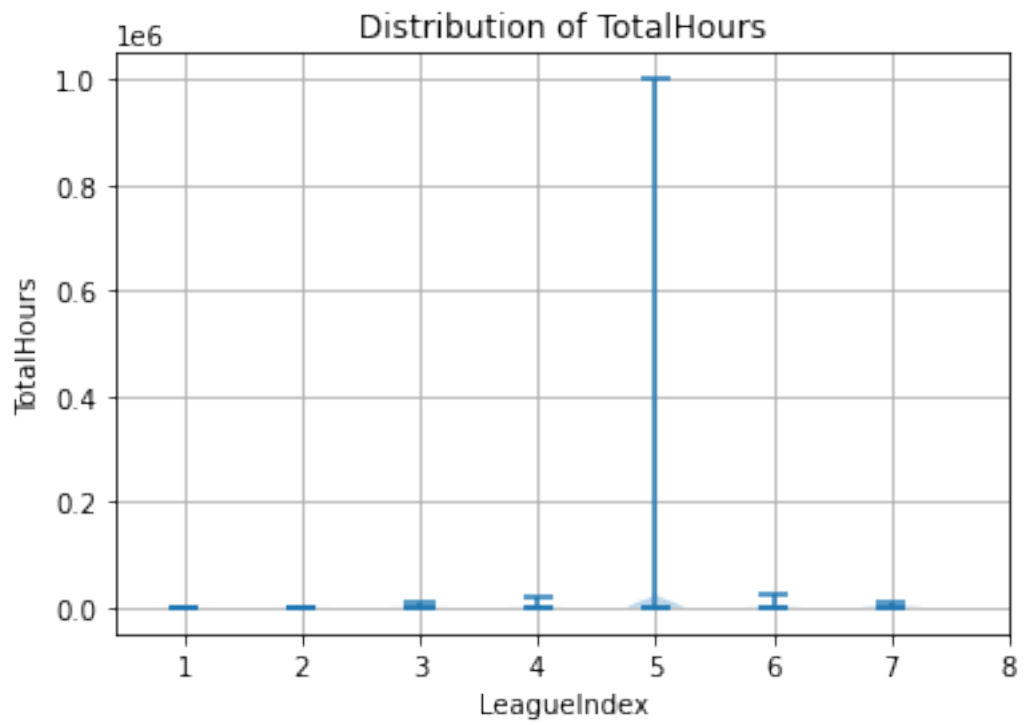
4 Visualizing Variables of Interest

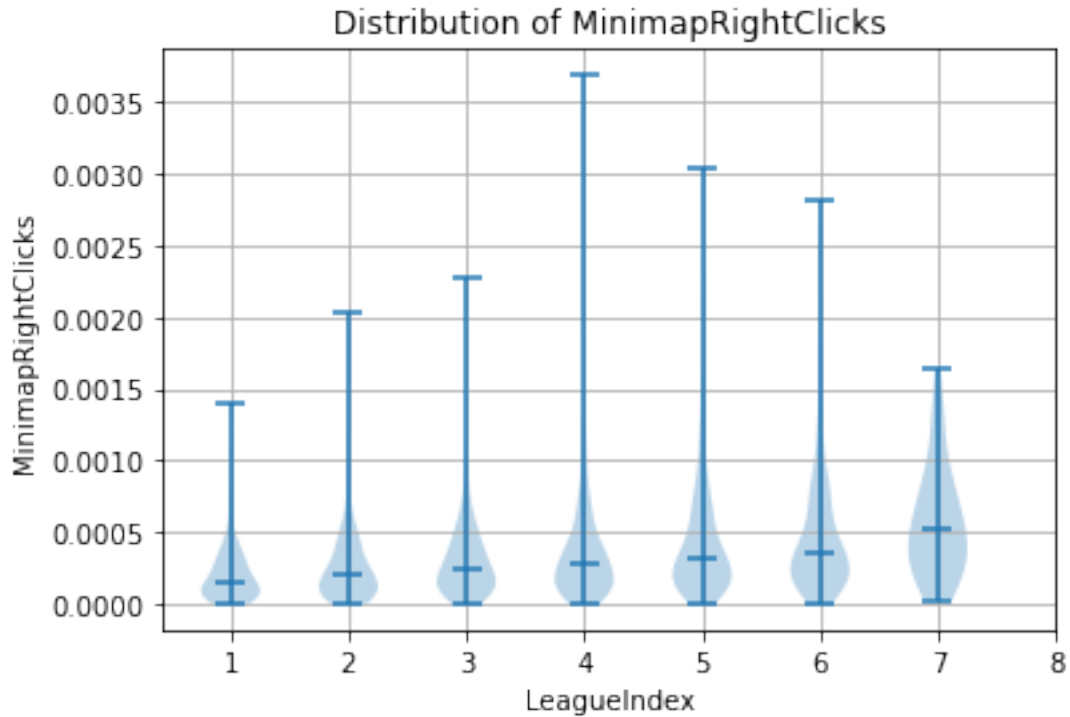
```
[57]: import matplotlib.pyplot as plt
```

```
[75]: var = ['APM', 'NumberOfPACs', 'GapBetweenPACs',
            'ActionLatency', 'TotalHours', 'SelectByHotkeys', 'MinimapRightClicks']
for v in var:
    data_list = []
    for i in range(1,8):
        data = df_without_question_mark[df_without_question_mark['LeagueIndex']_
↵== i][v]
        data_list.append(data)
    plt.violinplot(data_list, showmedians=True)
    plt.xlabel('LeagueIndex')
    plt.ylabel(v)
    plt.title(f'Distribution of {v}')
    plt.xticks(range(1, 9), range(1, 9))
    plt.grid(True)
    plt.show()
```



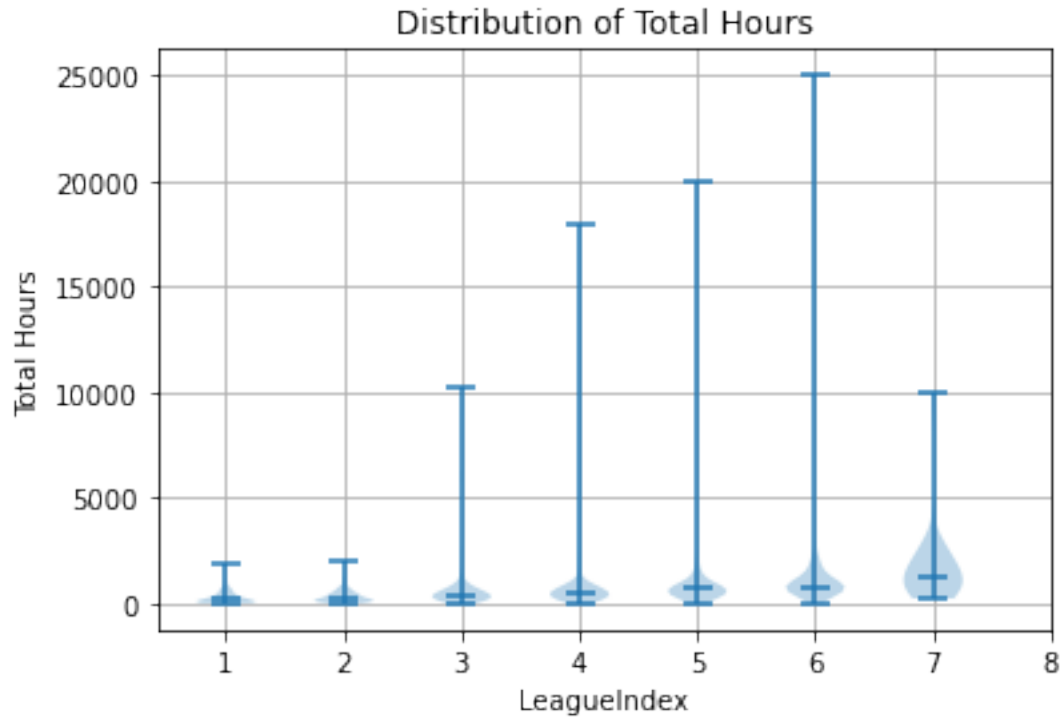






Looking at this we see an outlier in **TotalHours** that makes it hard to visualize so I'll remove that to see if there is a clear relation with our target variable

```
[80]: data_list = []
for i in range(1,8):
    data = df_without_question_mark[(df_without_question_mark['LeagueIndex'] == i) & (df_without_question_mark['TotalHours'] < 10**6)]['TotalHours']
    data_list.append(data)
plt.violinplot(data_list, showmedians=True)
plt.xlabel('LeagueIndex')
plt.ylabel('Total Hours')
plt.title(f'Distribution of Total Hours')
plt.xticks(range(1, 9), range(1, 9)) # Set x-axis tick labels to match LeagueIndex values
plt.grid(True)
plt.show()
```



Now some relation does appear to be present when you follow the medians, so I will train the model on the classification model on these features with that row removed.

5 Train/Test Split For Classification Model

```
[115]: data = df_without_question_mark[df_without_question_mark['TotalHours'] < 10**6]
```

```
[123]: from sklearn.model_selection import train_test_split
X = data[['APM', 'NumberOfPACs', 'GapBetweenPACs',
          'ActionLatency', 'TotalHours', 'SelectByHotkeys', 'MinimapRightClicks']]
y = data['LeagueIndex']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
                                                    random_state=42)
```

```
[124]: print("X_train shape:", X_train.shape)
print("X_test shape:", X_test.shape)
print("y_train shape:", y_train.shape)
print("y_test shape:", y_test.shape)
```

```
X_train shape: (2669, 7)
X_test shape: (668, 7)
y_train shape: (2669,)
y_test shape: (668,)
```

Based on the nature of the classification task (distinct integer categories) I will first try using a Decision Tree classifier. This is because it has good interpretability and may fit the data well, but it is possible to overfit, so I will also try a Random Forest, and see if that gives better performance on the test data. If overfitting does not appear to be a problem, the Decision Tree is preferable because of its interpretability. Otherwise, the Random Forest will be the solution to overfitting.

```
[125]: from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import mean_squared_error
model = DecisionTreeClassifier()
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
print("RMSE for DT:", rmse)
```

RMSE for DT: 1.3430905463031382

```
[126]: from sklearn.ensemble import RandomForestClassifier
model = RandomForestClassifier()
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
print("RMSE for RF:", rmse)
```

RMSE for RF: 1.124708877834722

It appears that the Random Forest gives better performance than the Decision Tree. However, because both are not great, and average being off by a bit more than 1 whole rank, I will try some other models.

```
[127]: from sklearn.linear_model import LogisticRegression
model = LogisticRegression()
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
print("RMSE for Logistic Regression:", rmse)
```

RMSE for Logistic Regression: 1.0943513103291655

```
/Users/ethan/miniconda3/envs/data/lib/python3.9/site-
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_i = _check_optimize_result(

```
[128]: from sklearn.svm import SVC
model = SVC()
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
print("RMSE for SVC:", rmse)
```

RMSE for SVC: 1.2118427597058525

```
[129]: from sklearn.ensemble import GradientBoostingClassifier
model = GradientBoostingClassifier()
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
print("RMSE for Gradient Boosting Classifier:", rmse)
```

RMSE for Gradient Boosting Classifier: 1.1153528482269206

5.0.1 To get a better sense of the different classification models I will test on multiple train-test splits and plot the RMSE's

```
[131]: group_labels = ['Decision Tree', 'Random Forest', 'Logistic Regression', 'SVC',
    ↪ 'Gradient Boosting']

dt_rmse = []
for i in range(50):
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↪ random_state=42)
    model = DecisionTreeClassifier()
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    mse = mean_squared_error(y_test, y_pred)
    rmse = np.sqrt(mse)
    dt_rmse.append(rmse)

rf_rmse = []
for i in range(50):
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↪ random_state=42)
    model = RandomForestClassifier()
    model.fit(X_train, y_train)
```

```

y_pred = model.predict(X_test)
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
rf_rmse.append(rmse)

lr_rmse = []
for i in range(50):
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    random_state=42)
    model = LogisticRegression()
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    mse = mean_squared_error(y_test, y_pred)
    rmse = np.sqrt(mse)
    lr_rmse.append(rmse)

svc_rmse = []
for i in range(50):
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    random_state=42)
    model = SVC()
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    mse = mean_squared_error(y_test, y_pred)
    rmse = np.sqrt(mse)
    svc_rmse.append(rmse)

gb_rmse = []
for i in range(50):
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    random_state=42)
    model = GradientBoostingClassifier()
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    mse = mean_squared_error(y_test, y_pred)
    rmse = np.sqrt(mse)
    gb_rmse.append(rmse)

```

/Users/ethan/miniconda3/envs/data/lib/python3.9/site-packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-

```
regression
  n_iter_i = _check_optimize_result(
/Users/ethan/miniconda3/envs/data/lib/python3.9/site-
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
regression
  n_iter_i = _check_optimize_result(
/Users/ethan/miniconda3/envs/data/lib/python3.9/site-
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
regression
  n_iter_i = _check_optimize_result(
/Users/ethan/miniconda3/envs/data/lib/python3.9/site-
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
regression
  n_iter_i = _check_optimize_result(
/Users/ethan/miniconda3/envs/data/lib/python3.9/site-
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
regression
  n_iter_i = _check_optimize_result(
/Users/ethan/miniconda3/envs/data/lib/python3.9/site-
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed
```

```
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/Users/ethan/miniconda3/envs/data/lib/python3.9/site-
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/Users/ethan/miniconda3/envs/data/lib/python3.9/site-
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/Users/ethan/miniconda3/envs/data/lib/python3.9/site-
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/Users/ethan/miniconda3/envs/data/lib/python3.9/site-
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
`n_iter_i = _check_optimize_result(
/Users/ethan/miniconda3/envs/data/lib/python3.9/site-
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.`

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
`n_iter_i = _check_optimize_result(
/Users/ethan/miniconda3/envs/data/lib/python3.9/site-
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.`

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
`n_iter_i = _check_optimize_result(
/Users/ethan/miniconda3/envs/data/lib/python3.9/site-
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.`

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
`n_iter_i = _check_optimize_result(
/Users/ethan/miniconda3/envs/data/lib/python3.9/site-
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.`

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/Users/ethan/miniconda3/envs/data/lib/python3.9/site-
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/Users/ethan/miniconda3/envs/data/lib/python3.9/site-
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/Users/ethan/miniconda3/envs/data/lib/python3.9/site-
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/Users/ethan/miniconda3/envs/data/lib/python3.9/site-
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/Users/ethan/miniconda3/envs/data/lib/python3.9/site-
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed
to converge (status=1):
```

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/Users/ethan/miniconda3/envs/data/lib/python3.9/site-
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed
to converge (status=1):
```

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/Users/ethan/miniconda3/envs/data/lib/python3.9/site-
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed
to converge (status=1):
```

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/Users/ethan/miniconda3/envs/data/lib/python3.9/site-
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed
to converge (status=1):
```

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/Users/ethan/miniconda3/envs/data/lib/python3.9/site-
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed
to converge (status=1):
```

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/Users/ethan/miniconda3/envs/data/lib/python3.9/site-
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/Users/ethan/miniconda3/envs/data/lib/python3.9/site-
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/Users/ethan/miniconda3/envs/data/lib/python3.9/site-
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/Users/ethan/miniconda3/envs/data/lib/python3.9/site-
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
/Users/ethan/miniconda3/envs/data/lib/python3.9/site-  
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed  
to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(  
/Users/ethan/miniconda3/envs/data/lib/python3.9/site-  
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed  
to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(  
/Users/ethan/miniconda3/envs/data/lib/python3.9/site-  
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed  
to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(  
/Users/ethan/miniconda3/envs/data/lib/python3.9/site-  
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed  
to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(  
/Users/ethan/miniconda3/envs/data/lib/python3.9/site-  
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed  
to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/Users/ethan/miniconda3/envs/data/lib/python3.9/site-
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/Users/ethan/miniconda3/envs/data/lib/python3.9/site-
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/Users/ethan/miniconda3/envs/data/lib/python3.9/site-
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/Users/ethan/miniconda3/envs/data/lib/python3.9/site-
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

```
https://scikit-learn.org/stable/modules/linear_model.html#logistic-
regression
n_iter_i = _check_optimize_result(
/Users/ethan/miniconda3/envs/data/lib/python3.9/site-
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-
regression
n_iter_i = _check_optimize_result(
/Users/ethan/miniconda3/envs/data/lib/python3.9/site-
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-
regression
n_iter_i = _check_optimize_result(
/Users/ethan/miniconda3/envs/data/lib/python3.9/site-
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-
regression
n_iter_i = _check_optimize_result(
/Users/ethan/miniconda3/envs/data/lib/python3.9/site-
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-
regression
n_iter_i = _check_optimize_result(
/Users/ethan/miniconda3/envs/data/lib/python3.9/site-
```

```
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/Users/ethan/miniconda3/envs/data/lib/python3.9/site-
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/Users/ethan/miniconda3/envs/data/lib/python3.9/site-
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/Users/ethan/miniconda3/envs/data/lib/python3.9/site-
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/Users/ethan/miniconda3/envs/data/lib/python3.9/site-
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```


Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/Users/ethan/miniconda3/envs/data/lib/python3.9/site-
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/Users/ethan/miniconda3/envs/data/lib/python3.9/site-
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/Users/ethan/miniconda3/envs/data/lib/python3.9/site-
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/Users/ethan/miniconda3/envs/data/lib/python3.9/site-
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```

regression
    n_iter_i = _check_optimize_result(
/Users/ethan/miniconda3/envs/data/lib/python3.9/site-
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear_model.html#logistic-](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

regression

```

    n_iter_i = _check_optimize_result(

```

```

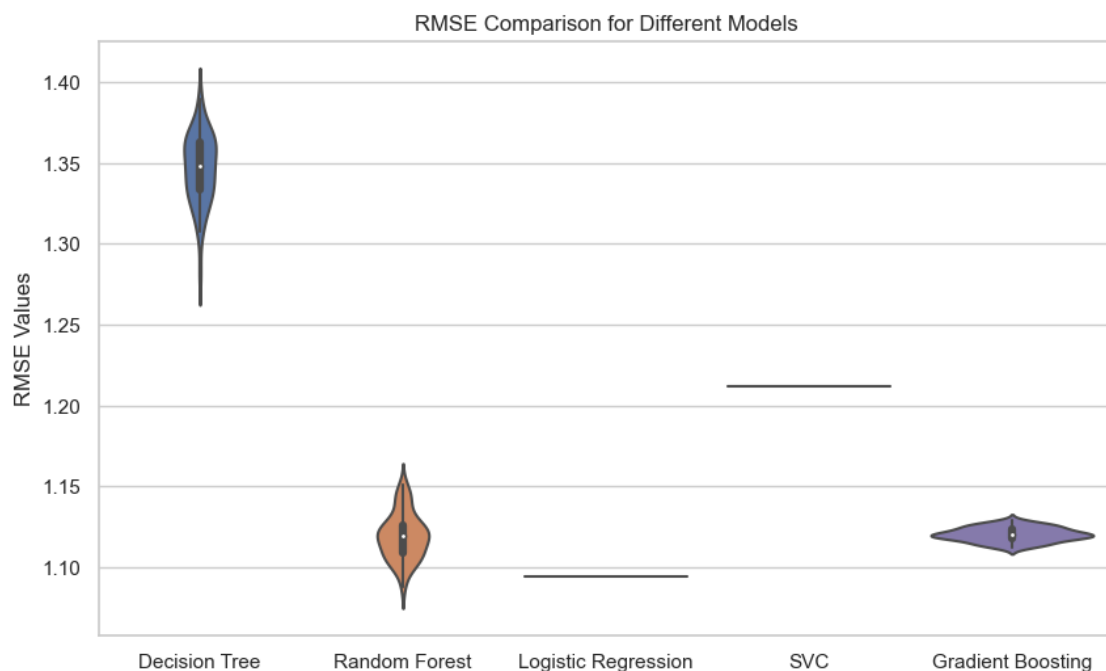
[135]: import seaborn as sns
import matplotlib.pyplot as plt
rmse_values = [dt_rmse, rf_rmse, lr_rmse, svc_rmse, gb_rmse]
sns.set(style="whitegrid")
fig, ax = plt.subplots(figsize=(10, 6), dpi=100)
ax = sns.violinplot(data=rmse_values)
ax.set_xticklabels(group_labels)
ax.set_ylabel("RMSE Values")
ax.set_title("RMSE Comparison for Different Models")

```

```

[135]: Text(0.5, 1.0, 'RMSE Comparison for Different Models')

```



6 Conclusion

It appears that the best model in terms of performance is the logistic regression. Although the performance is similar for many of the classification models, notably Logistic Regression performs the best and a Random Forest performs better than the Decision Tree classifier on the test set, so overfitting appears to not be an issue.

I would recommend the use of the logistic regression model, specified below because it has the lowest RMSE, and it also has very good interpretability, because it is easy to understand the logistic regression function and what the probability outputs mean.

```
[114]: model = LogisticRegression()
model.fit(X_train, y_train)
coefficients = model.coef_
intercept = model.intercept_
print("Coefficients:", coefficients)
print("Intercept:", intercept)
```

```
Coefficients: [[-3.97320020e-02 -1.15202002e-06  2.83030320e-02  3.42274734e-02
 -2.64148820e-03 -2.01214191e-06 -1.48514127e-07]
 [-1.67540557e-02 -5.11403476e-07  1.74990088e-02  2.79831344e-02
 -1.56508736e-03 -1.95864396e-06 -1.59310977e-08]
 [-5.12243720e-03 -3.11215896e-07  4.21422373e-03  2.23381862e-02
 -8.53615791e-05 -1.68413988e-06  2.44393605e-08]
 [ 6.84625077e-03  9.55242512e-08  7.18671839e-03  4.90093157e-03
  5.04184438e-04 -1.28843521e-06  7.99465771e-08]
 [ 1.93061163e-02  3.38182409e-07 -1.13087052e-02 -1.44690293e-02
  1.12683397e-03  1.19024052e-06  8.01249152e-08]
 [ 2.75528230e-02  1.79623176e-06 -2.27566303e-02 -3.93508327e-02
  1.27938236e-03  3.61855671e-06  1.08678378e-09]
 [ 7.90330492e-03 -2.55299036e-07 -2.31376475e-02 -3.56298635e-02
  1.38153637e-03  2.13456373e-06 -2.11524114e-08]]
Intercept: [ 1.35985174e-06  1.91768477e-04  1.21479083e-04  1.72614782e-04
 -1.41906968e-05 -1.18956177e-04 -3.54075320e-04]
```

```
/Users/ethan/miniconda3/envs/data/lib/python3.9/site-
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

6.1 For non-technical stake holders

By exploring the data, I found that several factors would be useful predictors of LeagueIndex i.e. rank. These factors were: **APM**, **NumberOfPACs**, **GapBetweenPACs**, **ActionLatency**, **TotalHours**, **SelectByHotkeys**, **MinimapRightClicks**. Most of these factors make sense because they relate to the speed of the player, and faster players probably have higher ranks, while others are good indicators of how much information a player is gathering or how much experience they have, which are both good ways to learn about the score. I chose to focus on this subset of factors to avoid over-fitting, which is where our model is trained so closely on training data, it struggles in the future to make accurate predictions. I ended up with a model that is both straight forward, as we can see its mathematical definition below, and in testing had a root mean squared error of only slightly greater than 1, which means we are on average roughly around 1 rank off, so it is a fairly good predictor. It also has no variability, which is better than the models with similar performance, as seen above.

Thus, the logistic regression model above is a suitable predictor of LeagueIndex.